

Schema Design

schema, n. — a representation of a plan or theory in the form of an outline or model.

Schemas

- ◉ Table Schema (i.e. relation schema)
 - What is the table called?
 - What columns does it have? What are their data types?
- ◉ Database Schema
 - What tables are in the database?
 - How are tables related?

Data Modeling

- ◉ How do we represent real world relationships and properties in our program?
 - ...in a way that makes writing the program easy
 - ...while remaining flexible for future changes
 - ...oh, it also has to be fast (enough).

Designing a Schema

- ◉ Analysis
 - What does my program need to output?
 - What data will I need to produce that output?
- ◉ Conceptual Design
 - Conceptual entities and their relationships
- ◉ Logical Design
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- ◉ Physical Design
 - JavaScript code, CREATE TABLE statements

Designing a Schema

What we'll focus on today

- Analysis
 - What does my program need to output?
 - What data will I need to produce that output?
- Conceptual Design
 - Conceptual entities and their relationships
- Logical Design
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- Physical Design
 - JavaScript code, CREATE TABLE statements

Designing a Schema

- Analysis
 - What does my program need to output?
 - What data will I need to produce that output?
- Conceptual Design
 - Conceptual entities and their relationships
- Logical Design
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- Physical Design
 - JavaScript code, CREATE TABLE statements

Example: A Journal Analysis


- ◉ I want a program to keep my journal in.
- ◉ I want to be able to enter the **text** of each journal entry.
- ◉ I want to be able to see journal entries **chronologically**.

Designing a Schema

- Analysis
 - What does my program need to output?
 - What data will I need to produce that output?
- Conceptual Design
 - Conceptual entities and their relationships
- Logical Design
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- Physical Design
 - JavaScript code, CREATE TABLE statements

Entity Relationship Diagram (ERD)

Conceptual Design



Designing a Schema

- Analysis
 - What does my program need to output?
 - What data will I need to produce that output?
- Conceptual Design
 - Conceptual entities and their relationships
- Logical Design
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- Physical Design
 - JavaScript code, CREATE TABLE statements

Entity Relationship Diagram (ERD)

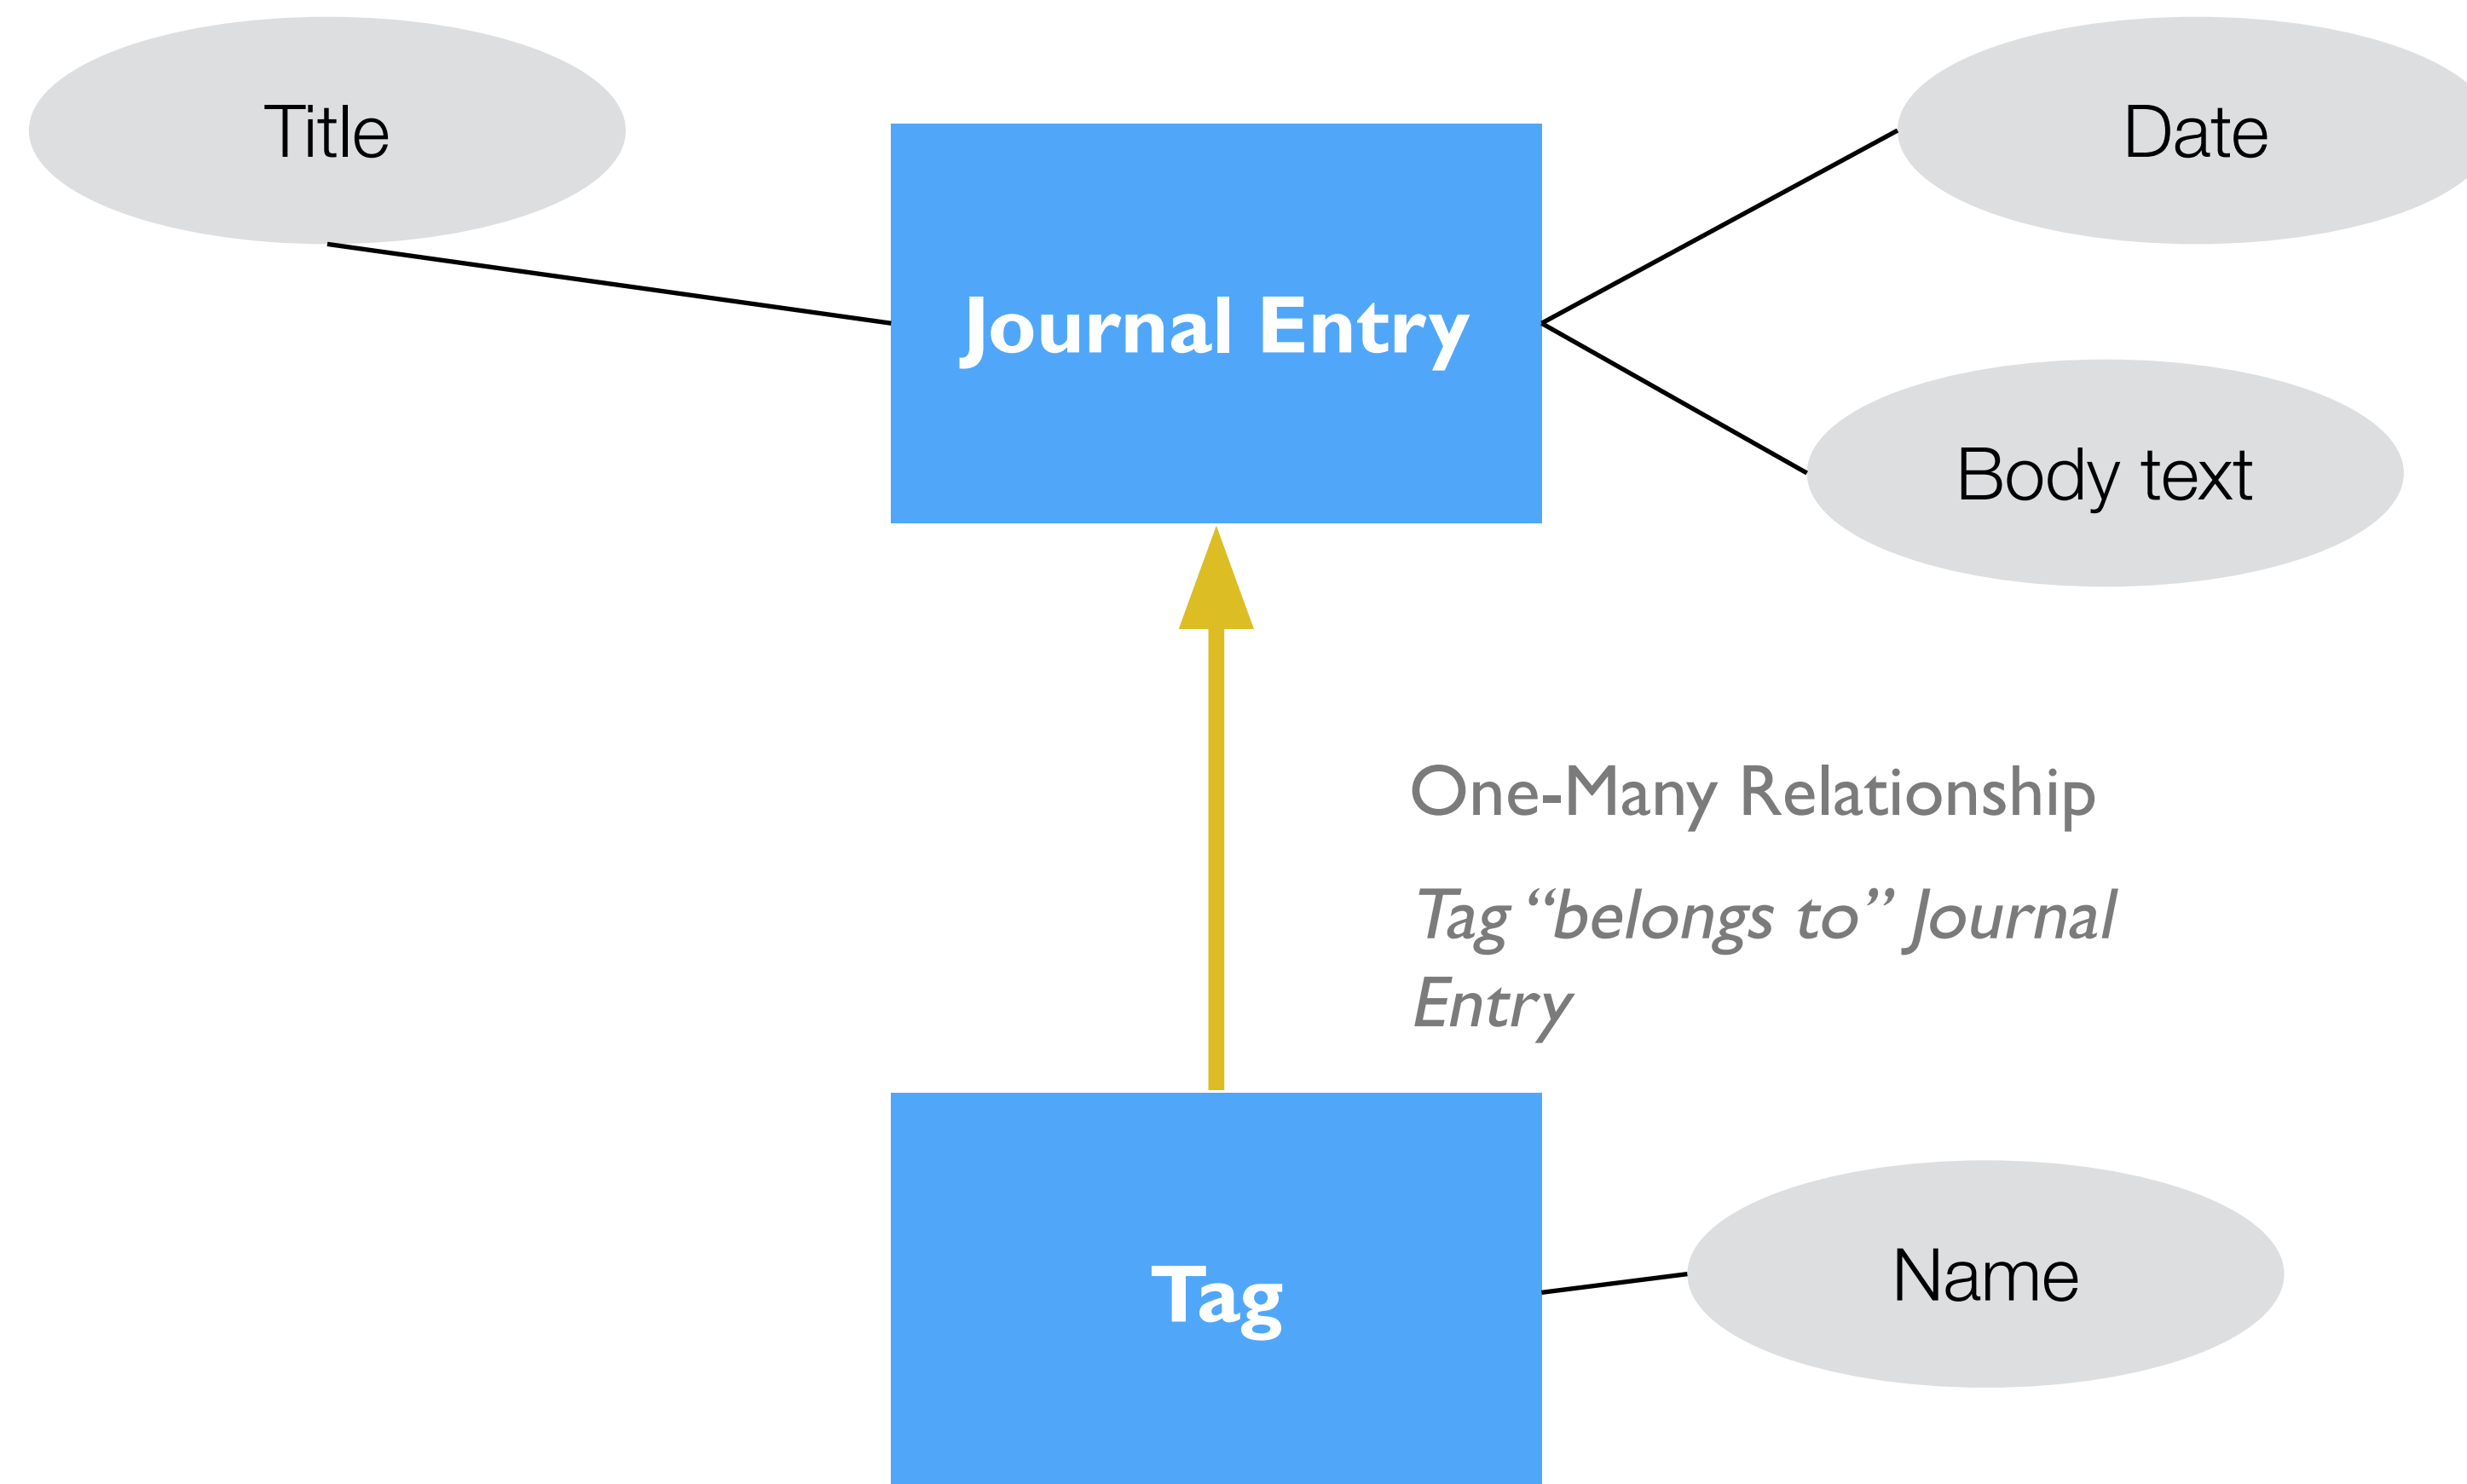
Logical Design

entries	
id	int, primary key
date_created	date
text	text

All done!

- ◉ Oh wait, I forgot a couple of things
 - I want to be able to have multiple journals
 - I want to be able to #tag entries and find all entries with a particular #tag
- ◉ Take 2...

Example: A Journal Conceptual Design, Take 2



Example: A Journal

Logical Design: Take 2

tagged_entries	
id	int, primary key
entry_id	int, foreign key
tag	string



entries	
id	int, primary key
date_created	date
text	text
journal_title	text

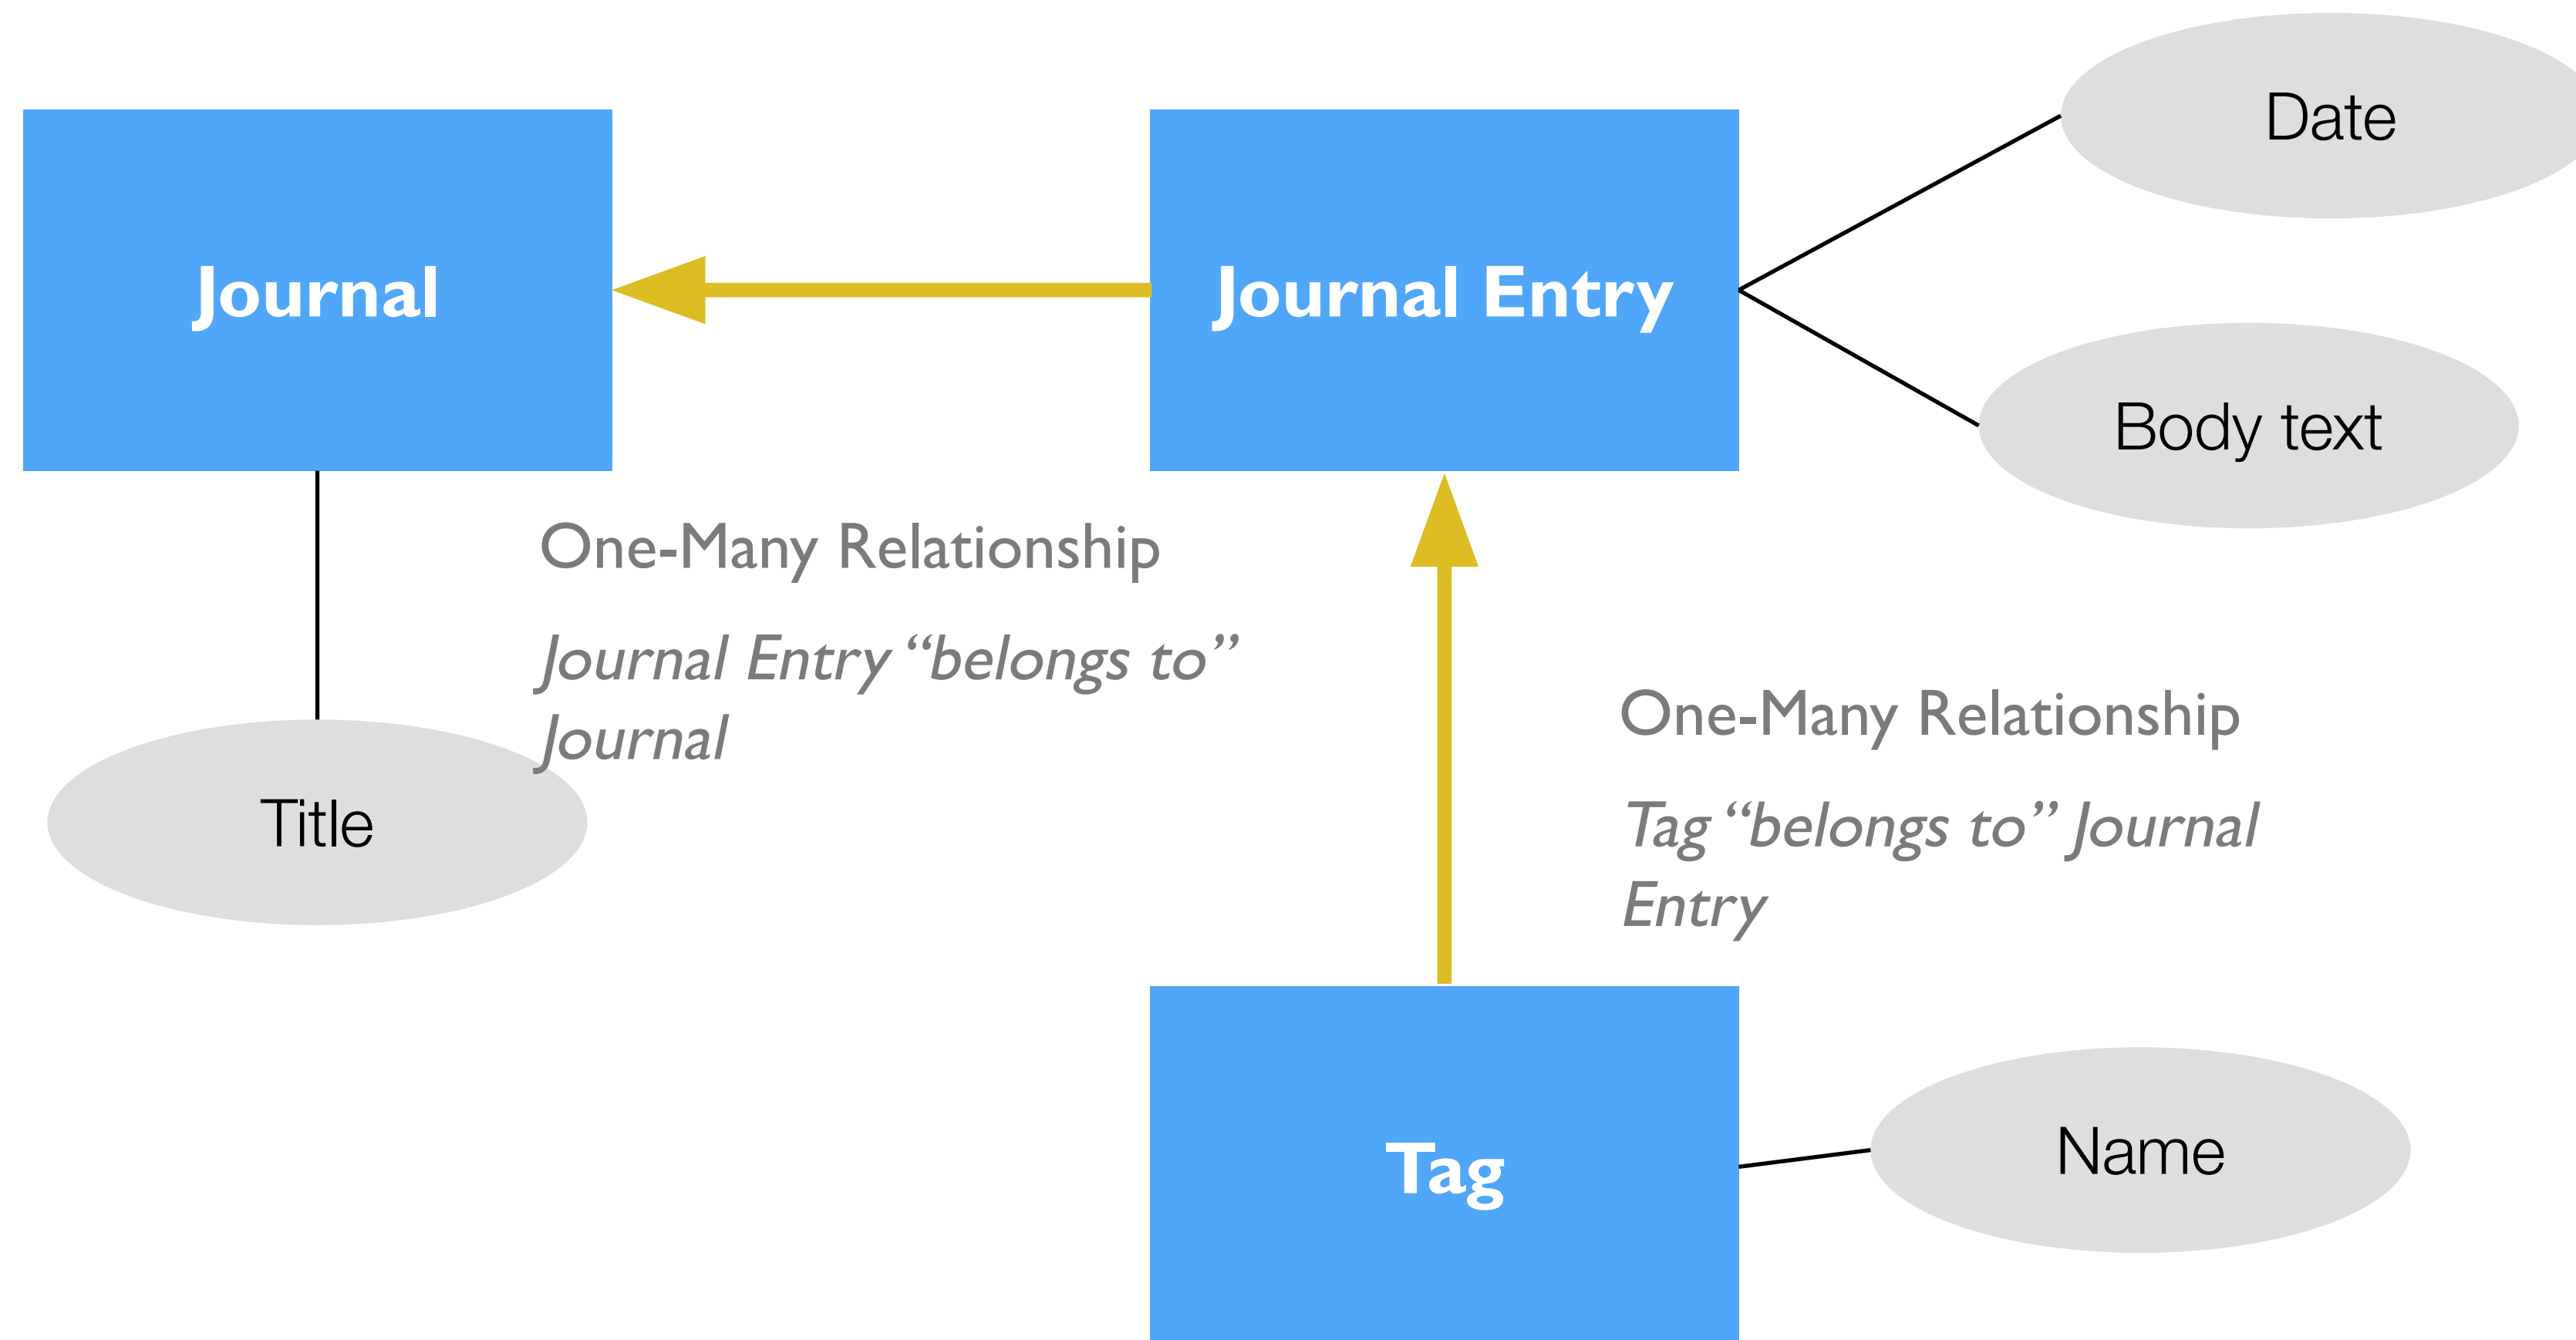
But Wait!!!

Normalization

- ◉ Organization that minimizes data redundancy and improves data integrity
- ◉ How do I change the name of “happy times” to “sadness”?


select * from entries;			
id	date_created	text	journal_title
0	2016-04-01	I am happy	happy times
1	2016-04-02	I am very happy	happy times
2	2016-04-03	Despair fills me	happy times
3	2016-04-03	Sadness is my life	an anatomy of pain

Conceptual Design, Take 3

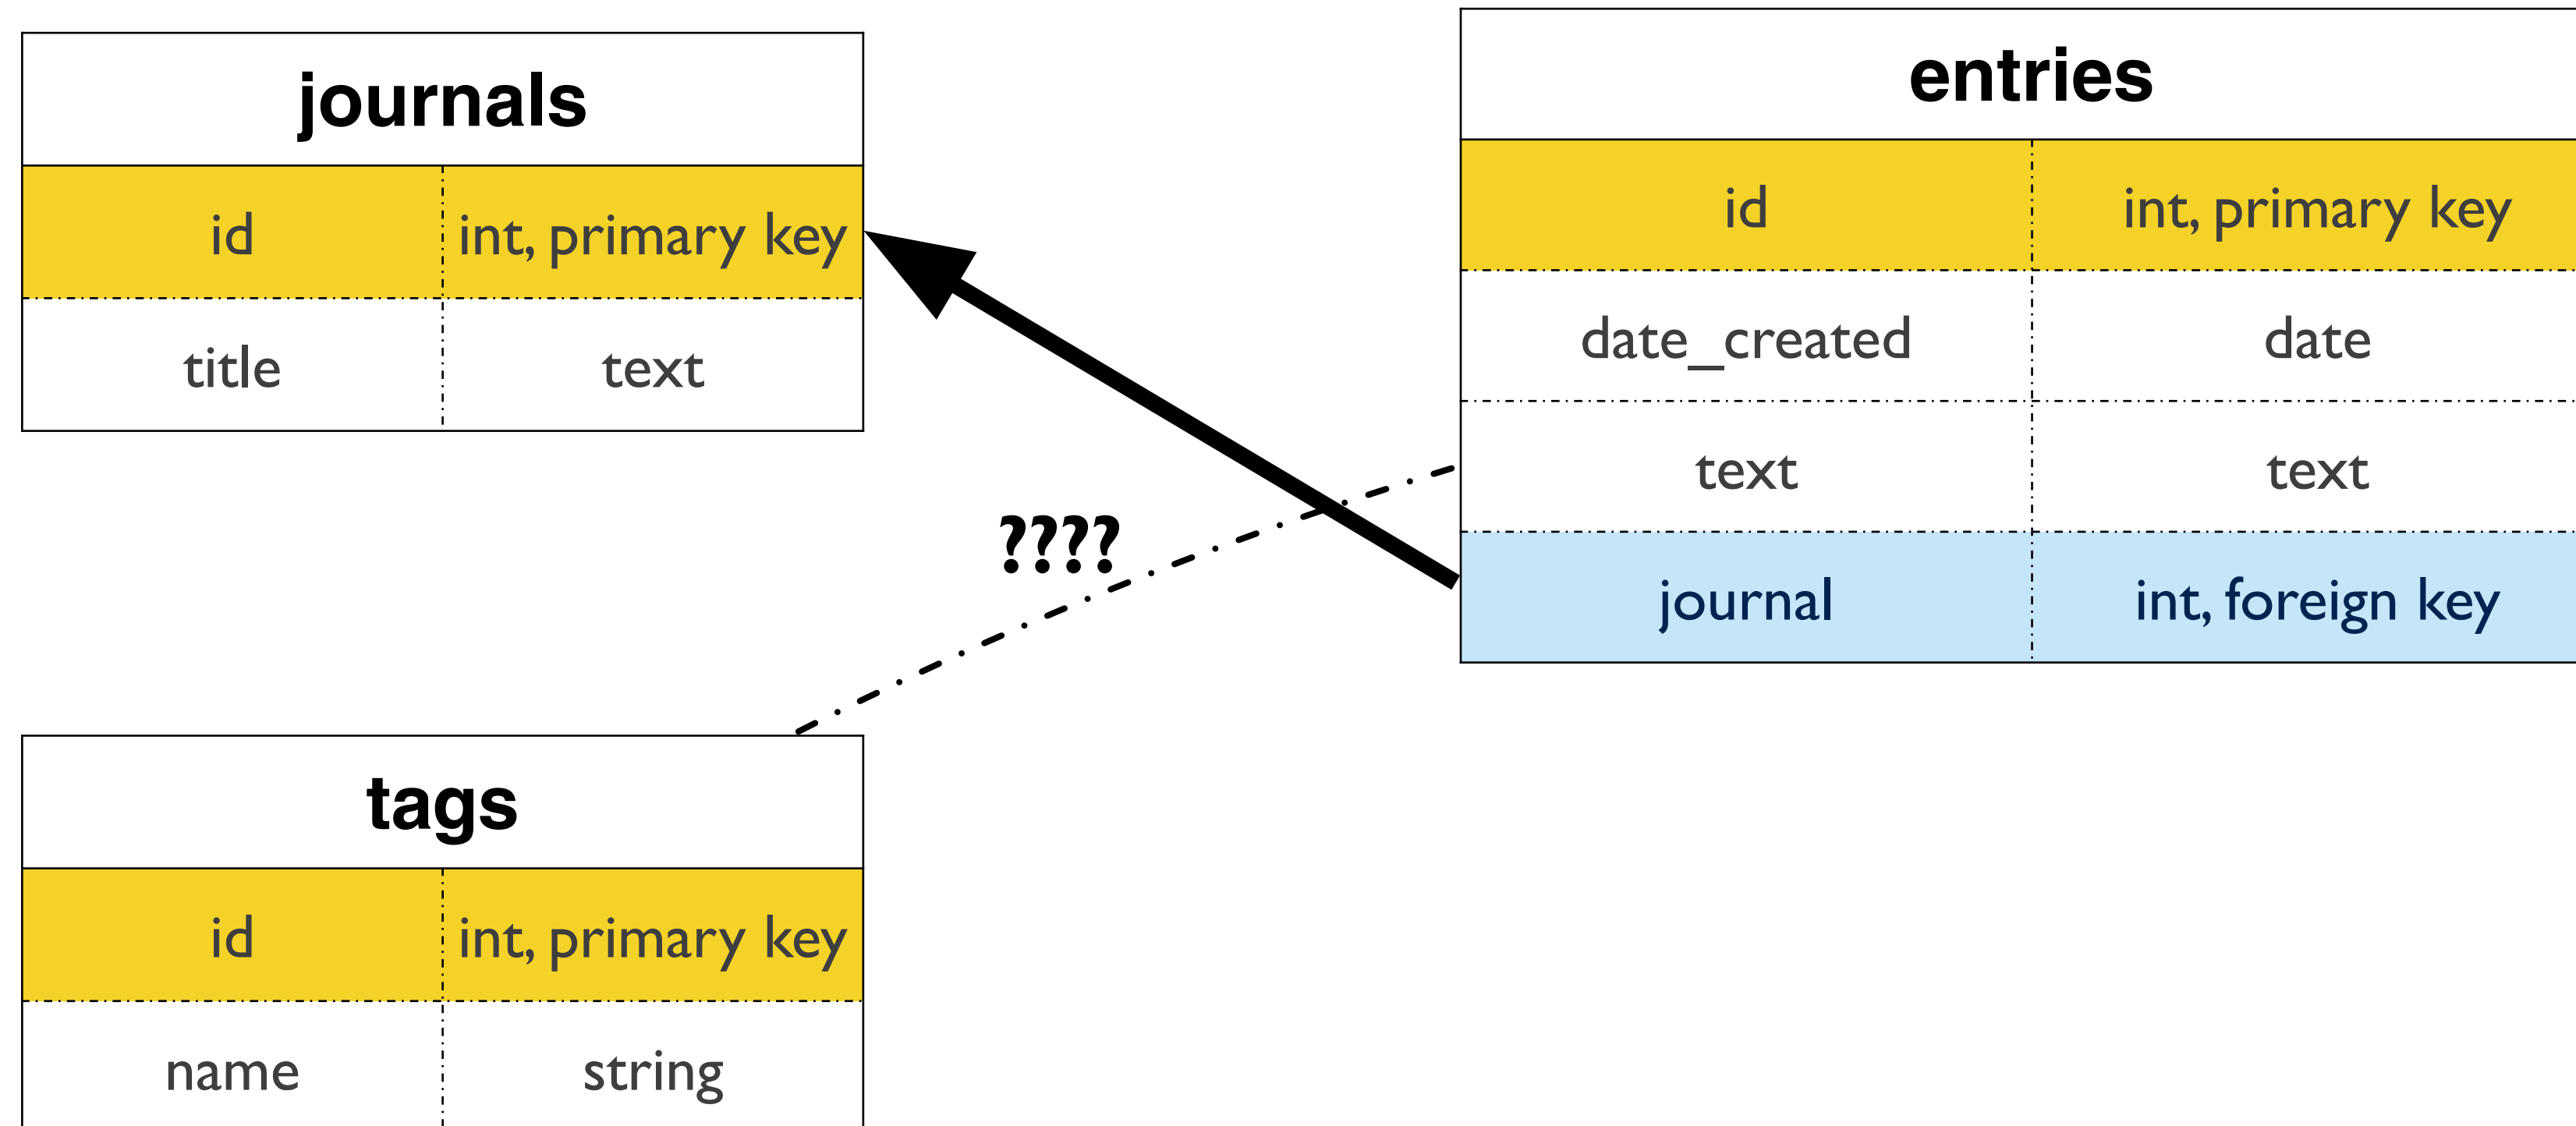


Logical Design, Take 3

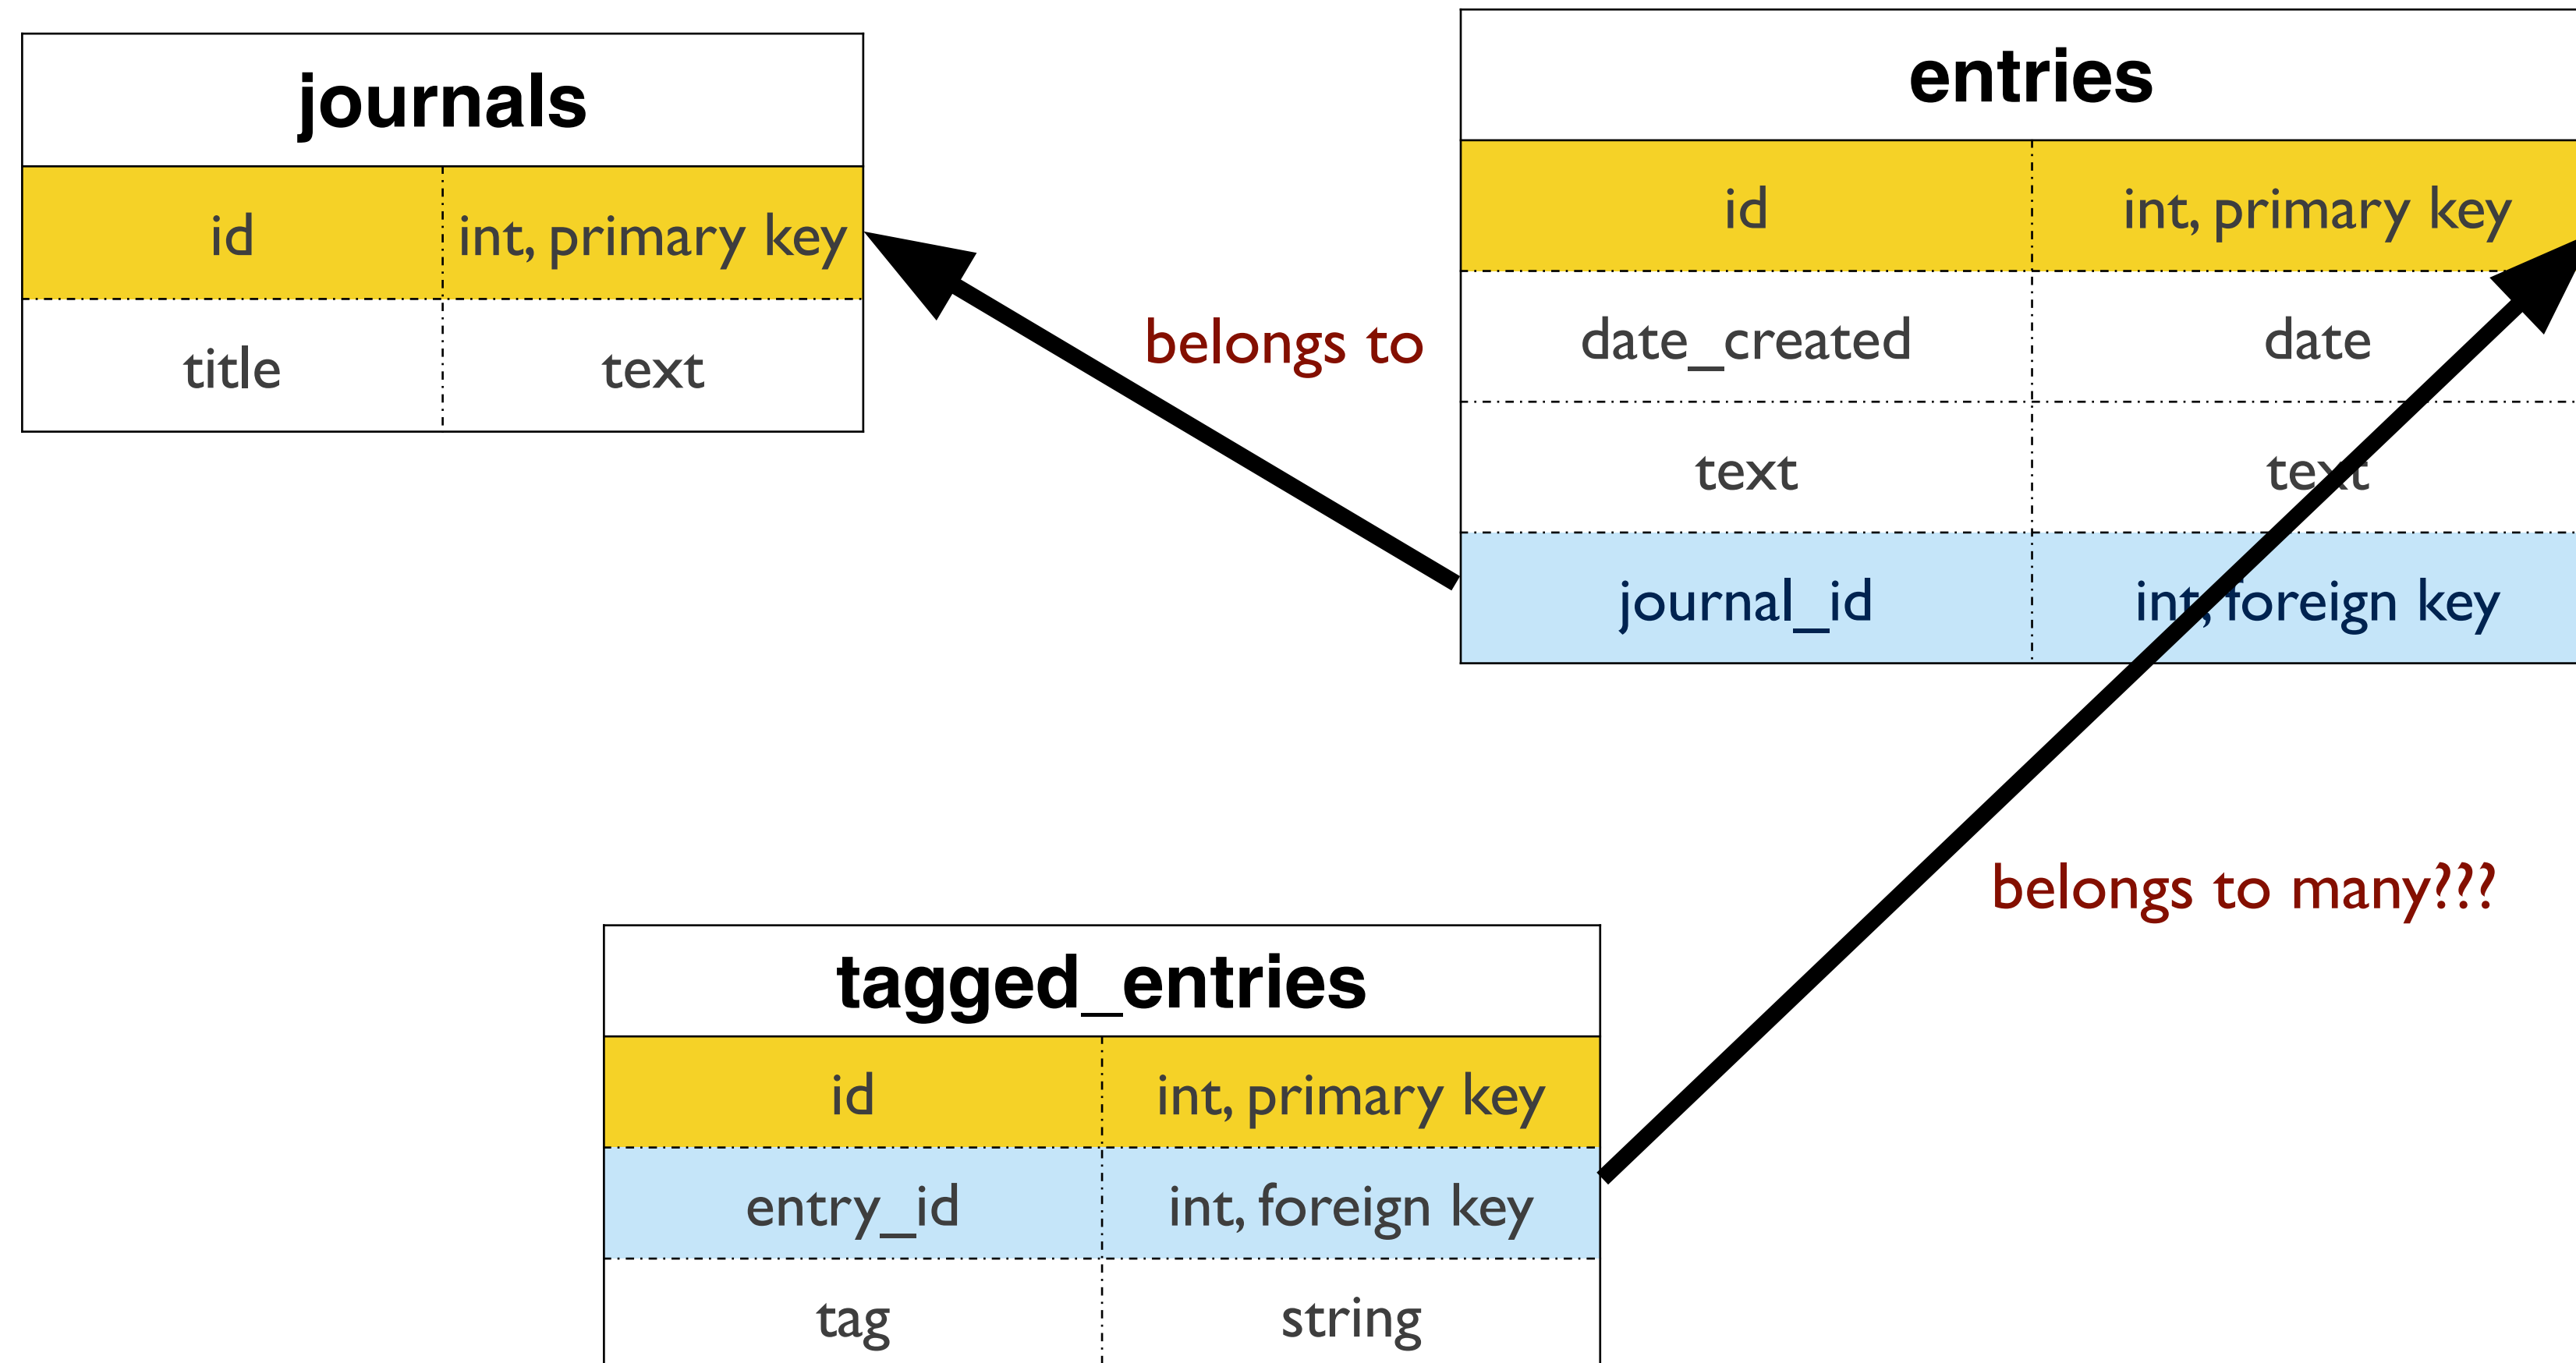
- Eliminate repeating groups in individual tables
- Create a separate table for each set of related data
- Identify each set of related data with a primary key



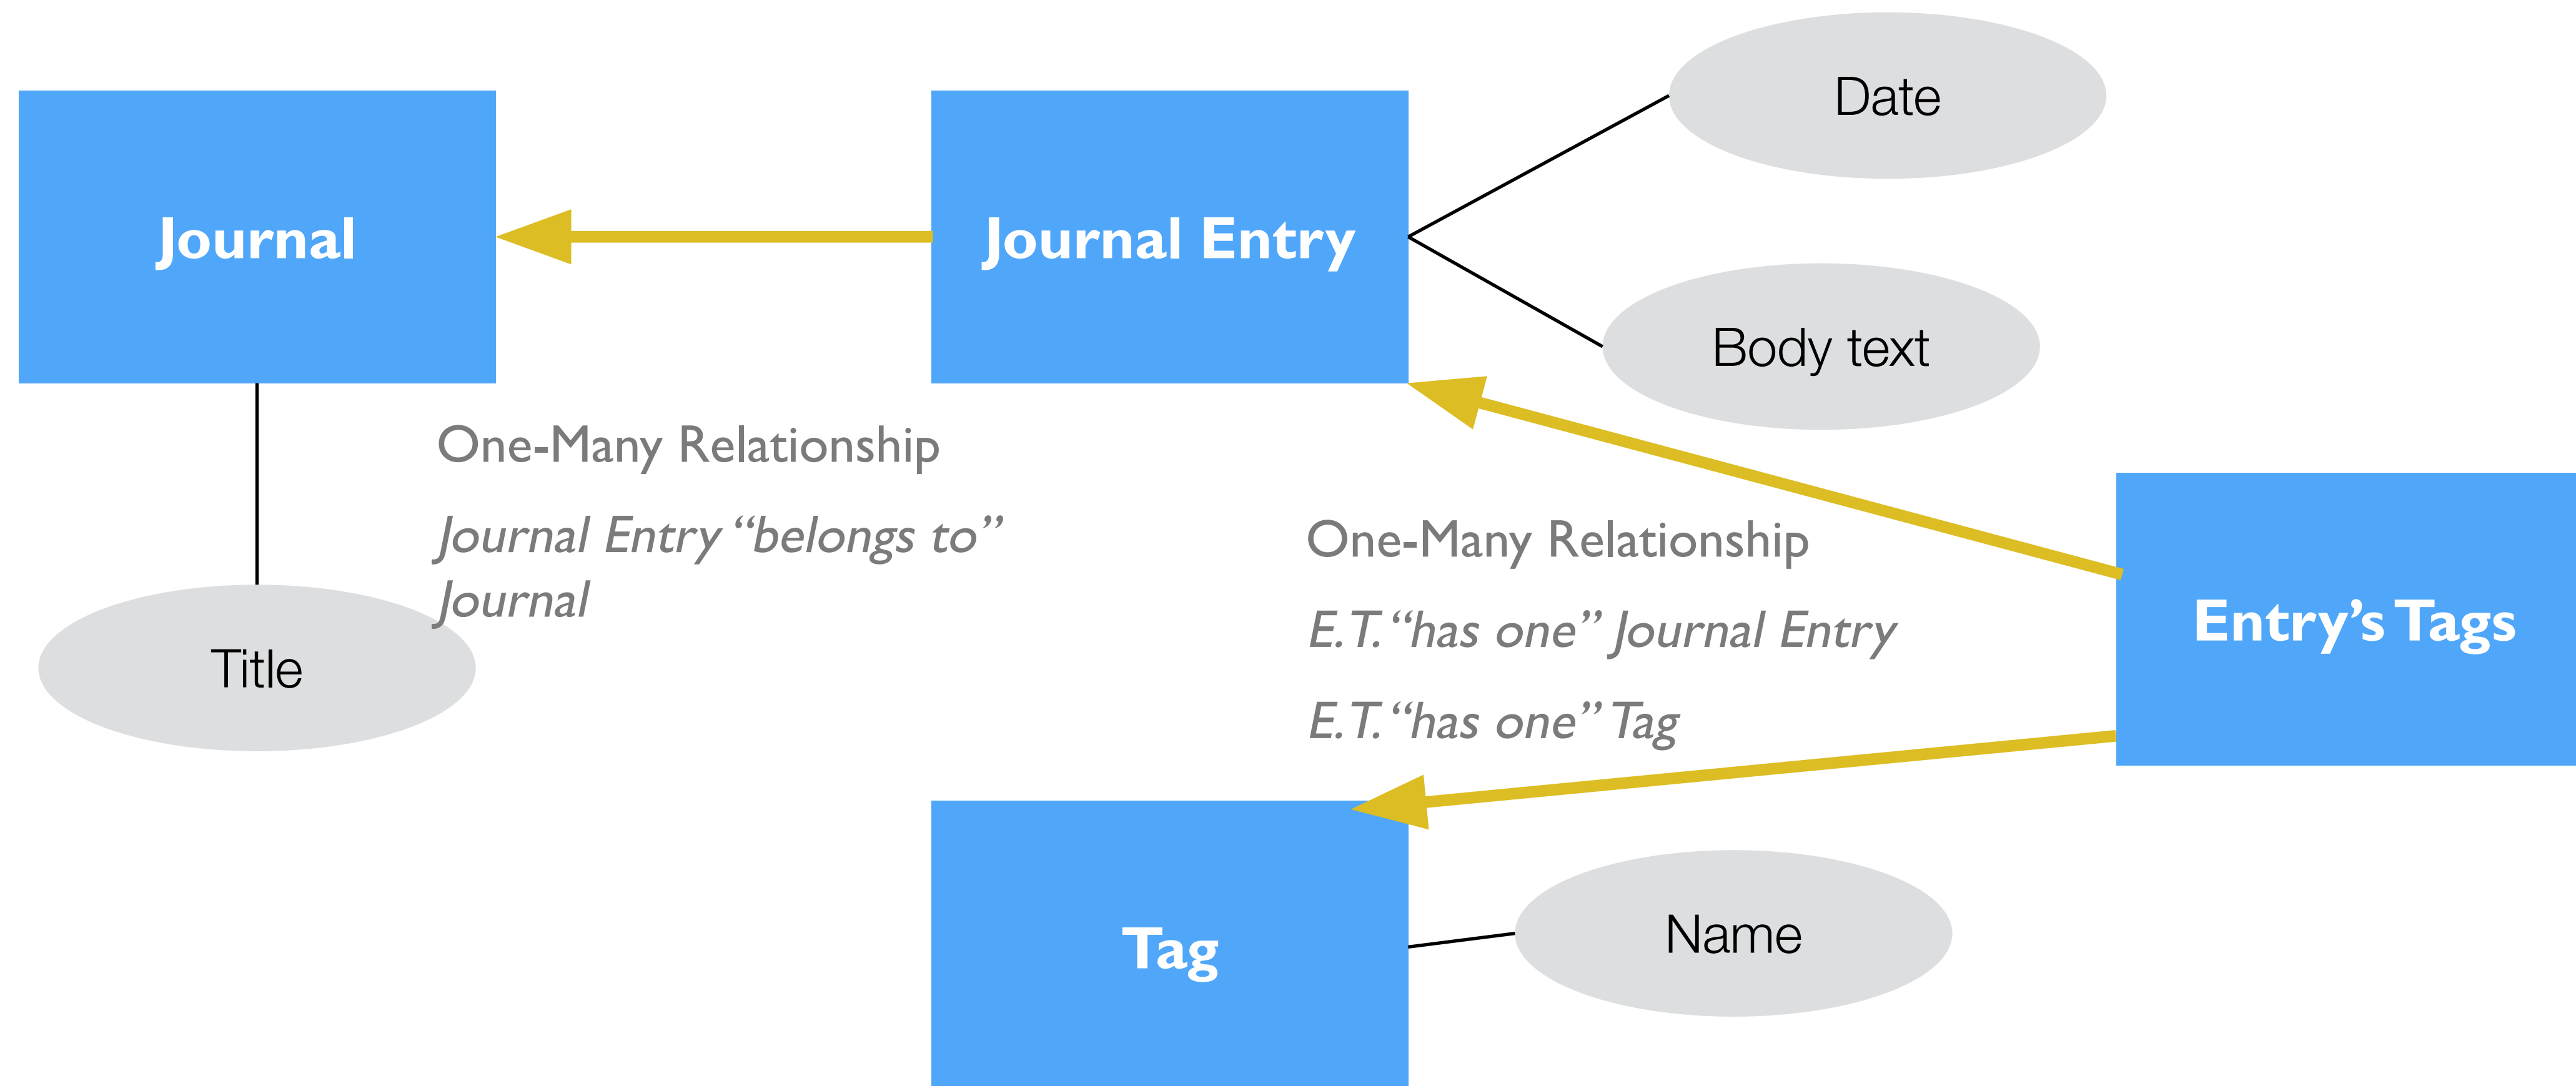
But what about tags?!?



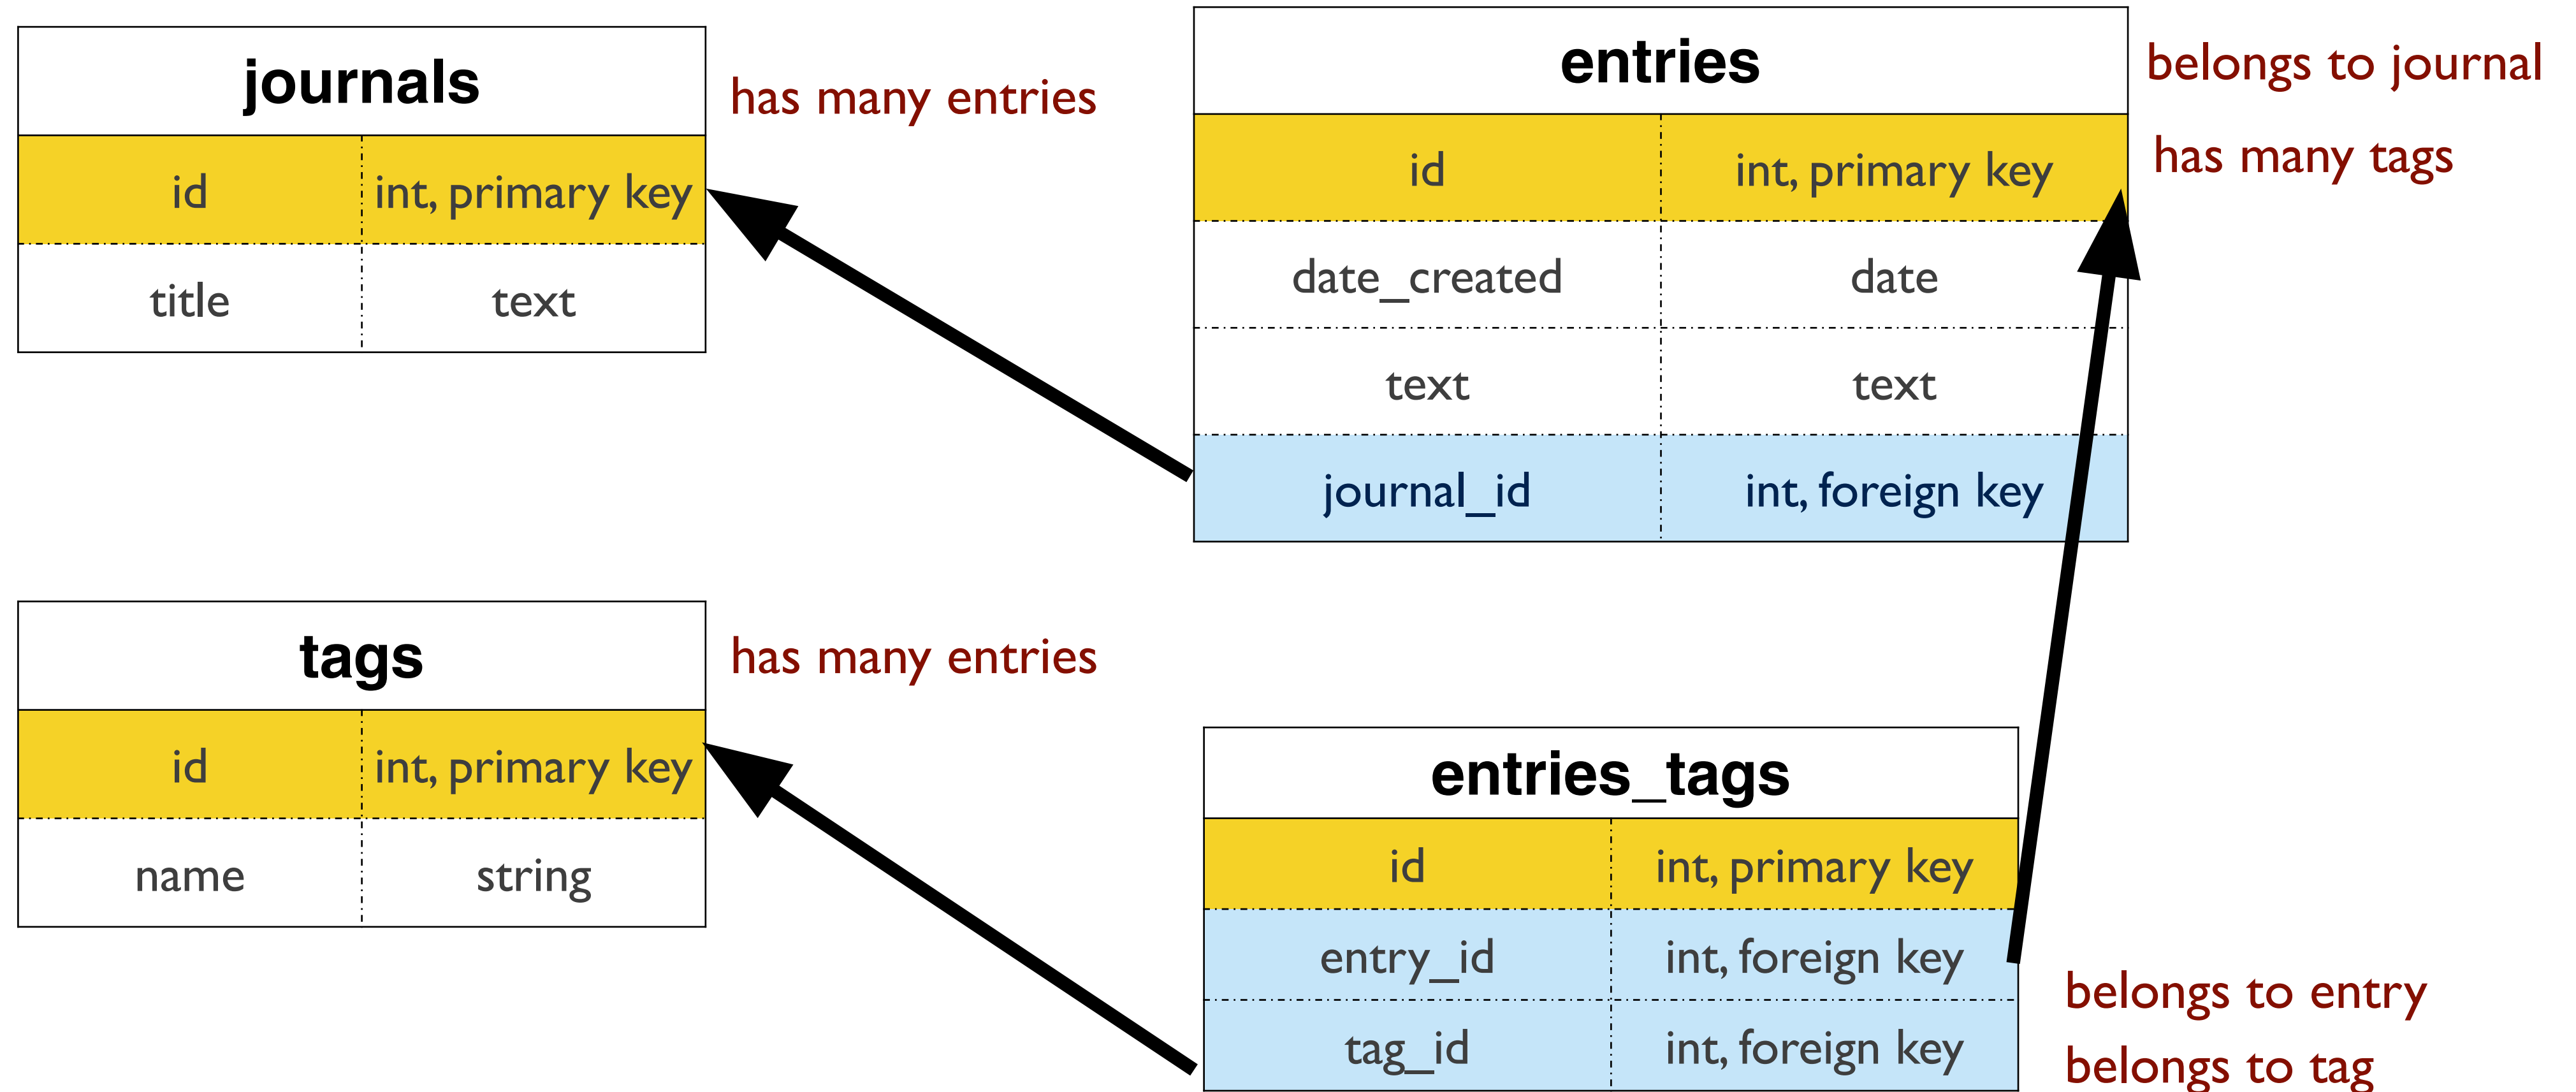
But what about tags?!?



Conceptual Design, Take 4



Logical Design, Take 4



Logical Design, Take 4

SELECT * FROM entries			
id	date_created	text	journal_id
0	2016-04-01	I am happy	0
1	2016-04-02	I am very happy	0
2	2016-04-03	Despair fills me	0
3	2016-04-03	Sadness is my life	1

SELECT * FROM journals		
id	date_created	title
0	2016-04-01	happy times
1	2016-04-02	an anatomy of pain

SELECT * FROM entries_tags	
tag_id	entry_id
0	0
0	1
1	0

SELECT * FROM tags		
id	date_created	name
0	2016-04-01	#YOLO
1	2016-04-02	#LOVELIFE

Has vs. Belongs To

- ◉ If rows in table A “belong to” rows in table B, that means A contains a foreign key for B
- ◉ If rows in table A “have” one or many rows in table B, that means table B is responsible for keeping track of the foreign key
- ◉ Think: the “owner” has less to worry about

Relationships

- ◉ Has One/Belongs To
 - Author has one Journal
 - Journal belongs to an Author
- ◉ Has Many/Belongs To
 - Entries belong to a Journal
 - A Journal has many Entries
- ◉ Belongs To Many
 - Tags and Journal Entries

Example: A Journal Normalized!

select * from entries;			
id	date_created	text	journal_id
0	2016-04-01	I am happy	0
1	2016-04-02	Very happy	0
2	2016-04-03	Despair...	0
3	2016-04-03	Sadness...	1

select * from entries_tags;		
id	entry_id	tag_id
0	0	0
1	0	1
2	1	0
3	1	2

select * from entries;	
id	name
0	0
1	0
2	1
3	1

select * from journals;	
id	title
0	happy times
1	an anatomy of pain

Do we even need this?

Composite Primary Keys

select * from entries;			
id	date_created	text	journal_id
0	2016-04-01	I am happy	0
1	2016-04-02	Very happy	0
2	2016-04-03	Despair...	0
3	2016-04-03	Sadness...	1

select * from entries_tags;	
entry_id	tag_id
0	0
0	1
1	0
1	2

select * from entries;	
id	name
0	0
1	0
2	1
3	1

select * from journals;	
id	title
0	happy times
1	an anatomy of pain

composite primary key

Normalized Databases

- ◉ Focus on optimal storage - often at odds with retrieval speed due to complex queries using complicated joins
- ◉ Work best when the application is write-intensive and write-load is more than read-load
 - Tables are usually smaller as data is divided vertically (fast reads on single tables)
 - Updates and Inserts are fast because there are no duplicates to update
 - Data is not duplicated so there is less of a need for process intensive group by or distinct queries
- ◉ Normalized tables mean join tables, which mean read operations on multiple tables suffer (indexing strategies don't work as well with joins)

Denormalized

- ◉ Works best when the application is read-intensive
 - The data is present in the same table (no need for joins)
 - A single table with all required data allows for efficient index usage
- ◉ Data is duplicated which means that updates and inserts become complex and costly

What Do I Do?!

- ◉ Real world applications will most likely have both read-loads and write-loads
- ◉ Utilize both approaches depending on the situation!
- ◉ Befriend your local DBA

Steps for Developing your ERD

1. Identify Entities
2. Define Relationships
3. Draw Rough-Draft ERD
4. Fill in Cardinality/Modality (arrows with relationship type)
5. Define Primary Keys
6. Label Foreign Keys
7. Identify and Map Attributes

Design Twitter!

